

ENGR-E 533 “Deep Learning Systems”

Lecture 03: Deep Learning Toolboxes

Minje Kim

Department of Intelligent Systems Engineering

Email: minje@indiana.edu

Website: <http://minjekim.com>

Research Group: <http://saige.sice.indiana.edu>

Meeting Request: <http://doodle.com/minje>



INDIANA UNIVERSITY

**SCHOOL OF INFORMATICS,
COMPUTING, AND ENGINEERING**

Intro to the GPU Resources

- A little bit about GPU computing
 - GPUs are different from CPUs in terms of their number of cores
 - Tesla V100 has 5120 FP32 cores
 - A better choice for a task with many simple sub-tasks, such as matrix multiplication
 - Each core is much less powerful than a usual CPU core
- Common packages you might need
 - Tensorflow or PyTorch
 - You may want to use some even higher lever wrappers like Keras, but I wouldn't cover them during the class
 - Jupyter
 - A nice interactive development framework
 - Tensorboard
 - If you want to see what's going on in your network



Baby TensorFlow

- Computational Graphs

- In TF things are represented with computational graphs
 - That consist of tensors and operations on them
- Building and running the computational graphs are separate
 - You define the graph first
 - This procedure is to represent what you want to do symbolically
 - It doesn't do any computation
 - Running the graph
 - Visit the graph nodes as required by the session's **run** method
 - Actually does the computation
- What for?
 - One obvious reason is to make your life easier by doing differentiation for you
 - So, it does the differentiation on the symbolic representations
 - And then actually calculate the gradient when you run the graph



Baby TensorFlow

- Tensors

o Three different kinds of tensors

- `tf.Variable`: something TF can change
- `tf.placeholder`: something you can change
- `tf.constant`: something nobody doesn't change once initialized

o For example,

- You want to train a classifier from a very large training dataset $\mathbf{Y} \approx \mathcal{G}(\mathbf{X}; \mathbb{W})$

This will be a very large data matrix

$$\frac{\partial \mathcal{E}}{\partial \mathbf{W}^{(2)}} = (\hat{\mathbf{Y}} - \mathbf{Y}) \mathbf{H}^T$$

A heavy matrix operation, maybe too large for the main memory

- So, you want to do something called **Stochastic** Gradient Descent (SGD), where
 - You sample a data point and calculate the gradient by using it
 - Or, you divide the entire data matrix into smaller pieces, called **minibatches** and calculate gradients for each of them
- Then,
 - `tf.Variable`: $\mathbf{W}^{(2)}$ ← You can set TF to update them during the run (of course you can initialize them as you want)
 - `tf.placeholder`: $\mathbf{H}_{:,t} \mathbf{Y}$ ← You need to feed this SGD samples or minibatches at every iteration
 - `tf.constant`: vector of ones for bias ← Once initialized nobody cares about this

Baby TensorFlow

- Running the graph

- If everything is symbolic, when do we actually do the operation?
 - In TF, there's a concept called “session” where all the variables and operations are actually taken care of
 - `tf.session.run(fetches, feed_dict=None, ...)`
 - `fetches`: The graph element you want to run
 - `feed_dict`: You can assign some values to the placeholder tensors
- Chain reaction
 - Running a graph element means running all the operations and evaluating tensors that are necessary for the fetched one
- For example,
 - You are interested in the training cost $\mathcal{E} = - \sum_t \sum_c Y_{c,t} \log \hat{Y}_{c,t}$
 - You define this cost as a graph element
 - If you do `sess.run()` on it, you need to know \hat{Y} , the prediction of the class labels
 - `sess.run()` traverses the graph to get this done
- `tf.gradients(cost, variables)`
 - Does the differentiation w.r.t. the variables
 - Most of the time you'll use a wrapper that calls this function internally



Baby PyTorch

- Difference between TF and PyTorch

○ The look

- TF: you construct the graph first and then run the graph
- PT: it does construct the graph but things are more seamless
 - Therefore the source code is more similar to the regular numpy codes

○ The tensors

- TF: there are three different types of tensors
- PT: a tensor is basically an N-dimensional array that has nothing to do with deep learning
 - But you can specify a tensor in the GPU memory if you want
 - `torch.autograd.Variable` defines a wrapper that turns a tensor into a PT variable
 - With which you can do all the cool things like automatic gradient calculation

○ GPU computing

- TF: if the same operation has two kernels for both CPU and GPU computing, GPU version gets the priority
- PT: there are some predefined GPU data types and a way to convert tensors
 - e.g. `torch.FloatTensor` versus `torch.cuda.FloatTensor`
 - `cuda()`: copies the tensor in the GPU memory



Baby PyTorch

- Difference between TF and PyTorch

○ The gradients

- TF: `tf.gradients(cost, variables)` differentiates the cost function w.r.t. the variable during the graph construction
 - The procedure is symbolic, so are the derivatives
- PT: all variables have their own `.grad` component that holds the actual gradient values
 - It's called "autograd" so things are still automatic, but treated differently
 - When `.backward()` method is called, the node is differentiated w.r.t. the leaf nodes in a recursive way
 - To do so, the forward pass records the input to the node
 - The backward pass calculate the gradients of all the intermediate nodes on the way back



TF versus PT

- TensorFlow

- Google
- Seems to have a larger user community
- Said to be better in building a serious project

- PyTorch

- Facebook
- Rapidly evolving
- Easier to quickly see the proof of concept

I do have my own preference, but for this course choose whatever you prefer



Thank You!



Minje Kim

Department of Intelligent Systems Engineering

Email: minje@indiana.edu

Website: <http://minjekim.com>

Research Group: <http://saige.sice.indiana.edu>

Meeting Request: <http://doodle.com/minje>



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING