

ENGR-E 533 “Deep Learning Systems”

Lecture 12: Generative Models

Minje Kim

Department of Intelligent Systems Engineering

Email: minje@indiana.edu

Website: <http://minjekim.com>

Research Group: <http://saige.sice.indiana.edu>

Meeting Request: <http://doodle.com/minje>





INDIANA UNIVERSITY

**SCHOOL OF INFORMATICS,
COMPUTING, AND ENGINEERING**

Latent Representation of Data

- A decoder made of deconvolution layers

- Suppose I want to create an image of a cat
 - What should the code vector look like?
 - Perhaps a one-hot vector?
- How about a dog?
- Any problem with this?



A colorful cat image

- We can't really distinguish different cats
- We can't really differentiate cats and lions versus cats and dogs

Latent Representation of Data

- Latent embedding vector

- How about this?

0	0.81	0	0	0	0	0	0	0	0	← A cat
0	0.78	0	0	0	0	0	0	0	0	← Another cat


- An even better representation

0	0.81	0	0	0	0.1	0	0	0	0	← A cat a little looking like a dog
0	0.78	0	0	0	0.55	0	0	0	0	← Another cat looking a lot like a dog

- And so on
- Eventually we need a latent representation with real numbers

Autoencoders with A Priori Knowledge

- Regularized autoencoders




- What do we want from the codes?
 - Minimal loss of information
- How do you define “information”?
 - The decoder should generate a similar output to the input
 - Originally similar input examples share similar codes
 - **As a generative model**
 - **Should be able to create examples from unseen codes**

Autoencoders with A Priori Knowledge


- Regularized autoencoders

- Problems with vanilla AEs
 - No control over the latent variables
 - Why do we need the control in the first place?
- A sample in the latent space
 - e.g. Could be something in between the two classes
 - Can we be more specific?
- The p.d.f. of the latent variables given the data set $p(\mathbf{z}|\mathbf{X})$
 - We want to know it better
 - This could be a nasty one
 - Furthermore, we want it to be more intuitively distributed
 - e.g. Z_1 is for thickness; Z_2 is for rotation; etc.
 - Why?
 - How could you “generate” a new example with a choice of thickness if the p.d.f. is too complicated?



Variational AutoEncoders

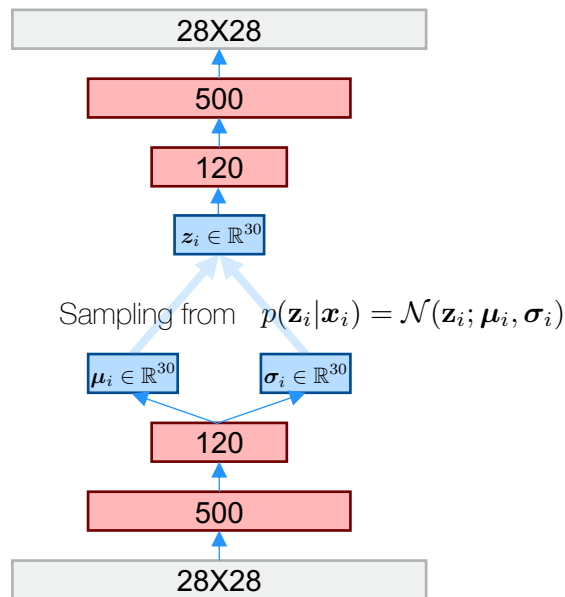
- Latent variables are from multivariate Normal distributions



Vanilla AE

Encoder

Variational AE




- Why this weird structure?

- You want an easy distribution for $p(\mathbf{z}_i | \mathbf{x}_i)$
- But it makes training difficult—needs a special structure

Variational AutoEncoders

- Latent variables are from multivariate Normal distributions

Sampling turns decoding into a stochastic process, enabling it to learn from multiple samples of the LV



Parameterization of the Normal distribution
lets encoding result in a distribution over LV

- During training
 - Every epoch the reconstructed input (output of the decoder) is from a sampled version of the latent variable
 - The training algorithm is exposed to the continuous distribution of the LV, which is not possible with vanilla AE
- During testing
 - The decoder itself is deterministic, but produces stochastic predictions due to the sampling process